Multi-Label Movie Genre Classification

Shashank V. Maiya

Springboard

Data Science Intensive Capstone Project

October 1, 2018 Cohort

Types of Classification Tasks

Binary Classification

- Spam
- Not spam

Binary Classification:
Only 2 choices

Multiclass Classification

- Dog
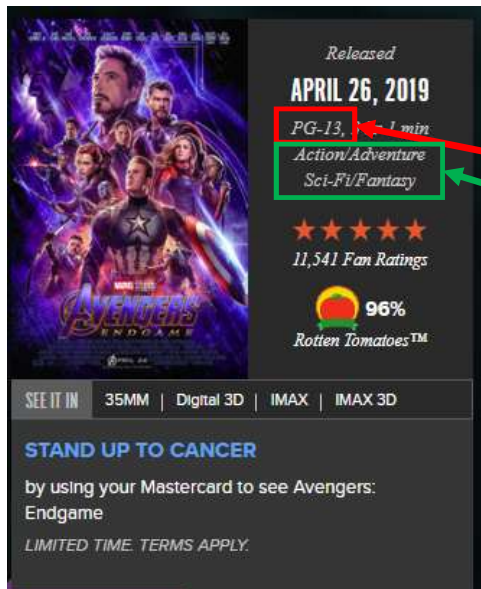- Cat
- Horse
- Fish
- Bird
- ...

Multi-class Classification:
Each observation is assigned to one and only one class

Multi-label Classification

- Dog
- Cat
- Horse
- Fish
- Bird
- ...

Multi-label Classification:
Each observation can be classified into multiple classes

Released
APRIL 26, 2019
PG-13, _ _ 1 min
Action/Adventure
Sci-Fi/Fantasy

★★★★★
11,541 Fan Ratings

96%
Rotten Tomatoes™

SEE IT IN   35MM | Digital 3D | IMAX | IMAX 3D

STAND UP TO CANCER
by using your Mastercard to see Avengers: Endgame
LIMITED TIME. TERMS APPLY.

Multi-Label Classification - Applications

Multi-Label Classification for Image/Scene Categorization

Tags:
- Lake
- Human
- Sky
- Hills
- Clouds

Horror: 0.02%
Romance: 0.02%
Adventure: 99.96%
Documentary: 0.0%

- Multi-Label Classification of Genres based on movie Posters

Multi-Label Classification of Movie Genre using Plot

Doctor Strange (2016)   7.5/10 424,942   Rate This
PG-13 | 1h 55min | Action, Adventure, Fantasy | 4 November 2016 (USA)

1:01 | Trailer   21 VIDEOS | 251 IMAGES

While on a journey of physical and spiritual healing, a brilliant neurosurgeon is drawn into the world of the mystic arts.

Prediction Problem

- Given the plot of the movie, what are the genres they fall into?

- Total of 27 Possible Genres

- Movies are classified anywhere from 1 to 12 genres

Who might care?

Online Streaming Companies

Movie Review Websites

Data Overview

- Data set obtained from IMDB

- Number of movie data = 117,194

- Column Description
  - Movie Title
  - Plot Language
  - Movie Plot
  - 27 Movie Genres

	title	plot	Action	Adult	Adventure	Animation	...	Sport	Talk-Show	Thriller	War	Western	plot_lang
0	"#7DaysLater" (2013)	dayslater interactive comedy series feature en...	0	0	0	0	...	0	0	0	0	0	en
1	"#BlackLove" (2015) {Crash the Party (#1.9)}	week leave workshops women consider idea ladie...	0	0	0	0	...	0	0	0	0	0	en

2 rows × 30 columns
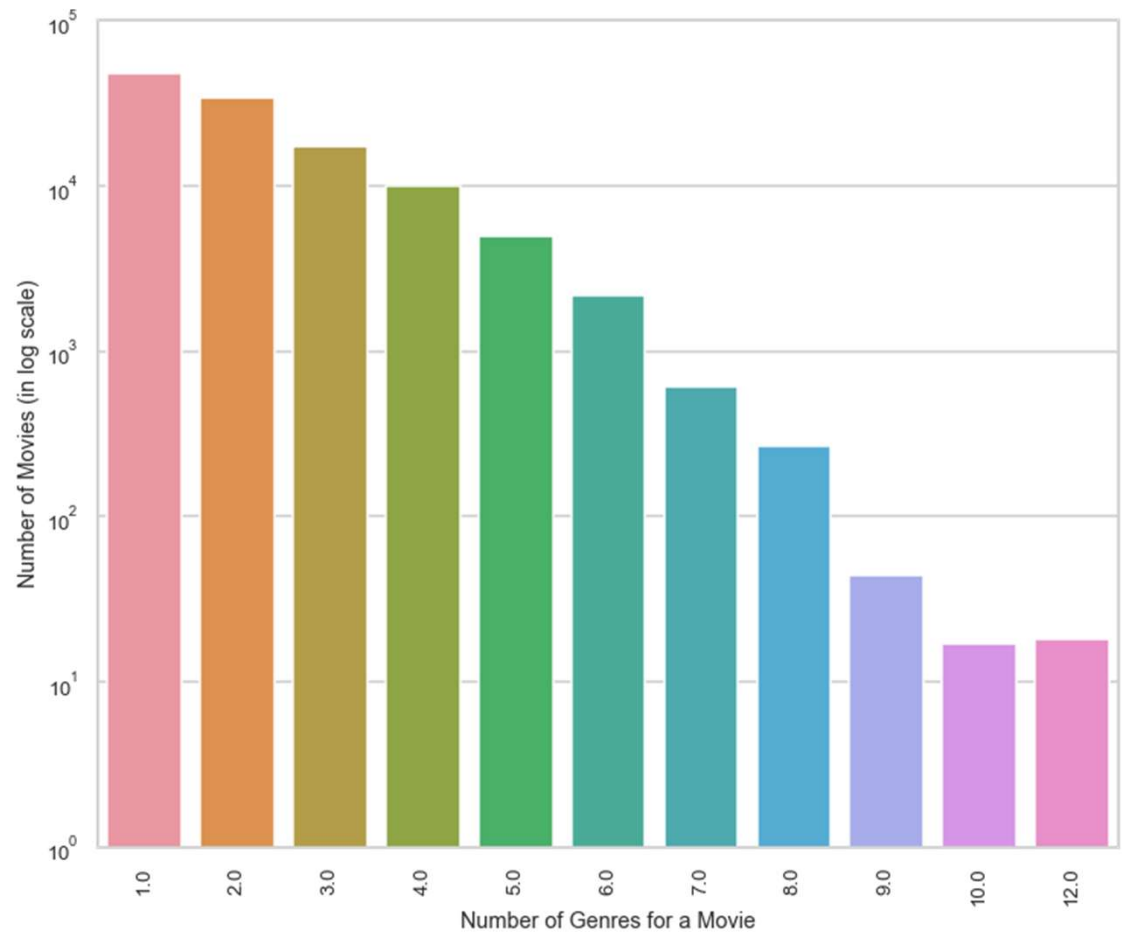
Exploratory Data Analysis (EDA)

Number of Movies per Genre

- Highest Genre Movies = Drama (45891) followed by Comedy (33870)
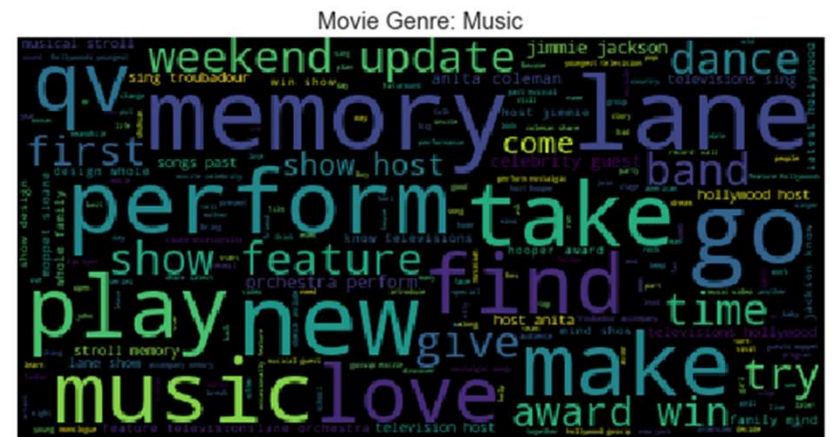- Lowest Genre Movies = Adult (61)

Number of Genres per Movie

- Most movies are categorized with 1 or 2 genres

- Average of 2.1 genres per movie

- 18 movies are classified using 12 genres!

Word Cloud plots – [Sports, Music]

Movie Genre: Sport
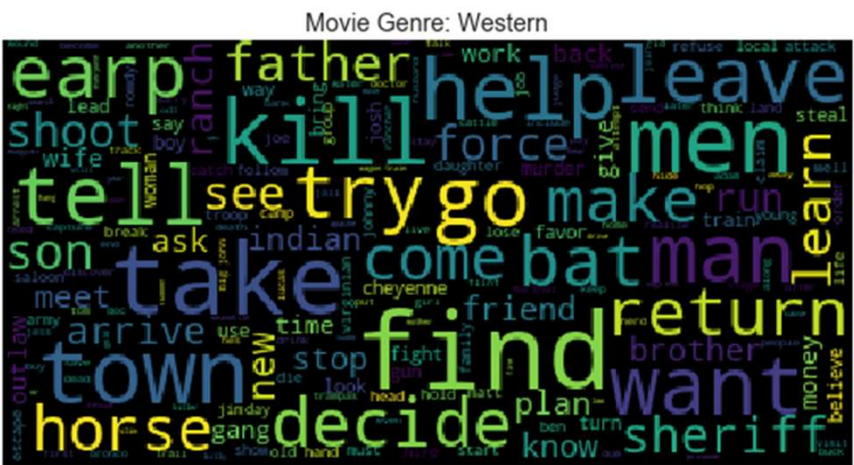
Relevant words – 'team, vs,
sport, first, match, player, race'

Movie Genre: Music

Relevant words – 'memory, lane,
perform, play, love, award, show'

Word Cloud plots – [Biography, War]

Movie Genre: Biography

Relevant words – 'life, career,
world, story, interview'

Movie Genre: War

Relevant words – 'german,
order, plan, kill, mission,
american'

Word Cloud plots – [Crime, Western]

Movie Genre: Crime

Relevant words – 'kill, murder, case, police, victim'

Movie Genre: Western

Relevant words – 'horse, sheriff, town, kill, Indian, gang'

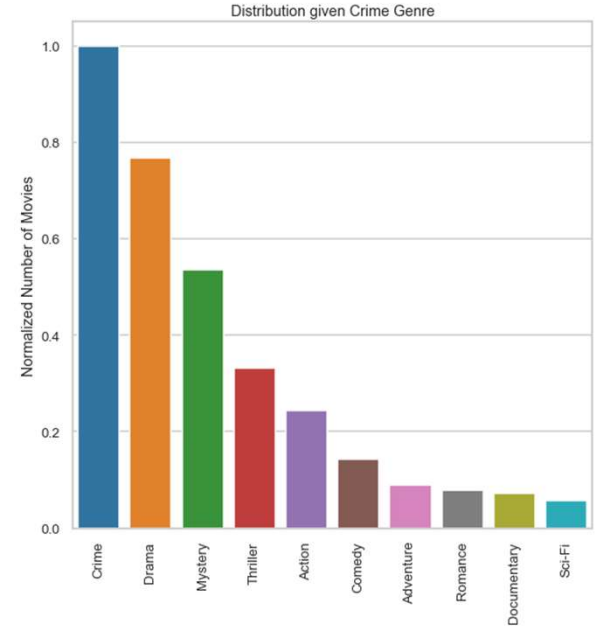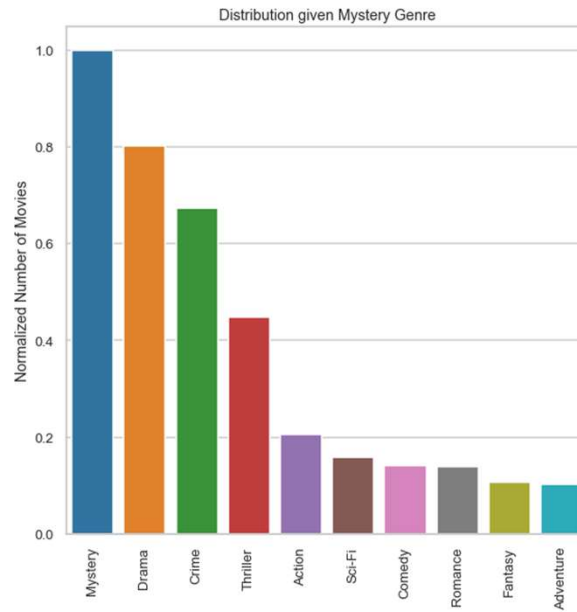# Correlation Analysis - Heatmap

Genres with strong positive correlation

- Action, Adventure, Fantasy & Sci-Fi
- Animation, Fantasy & Family
- Crime, Thriller, Mystery & Drama
- Drama & Romance
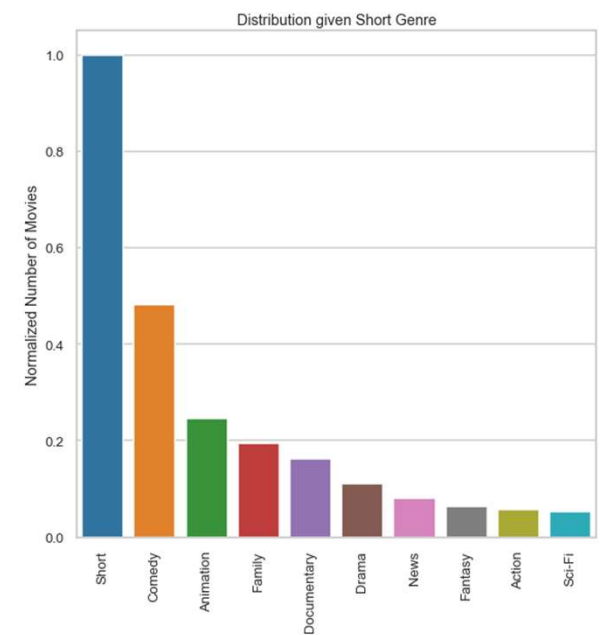- Game-Show & Reality-TV
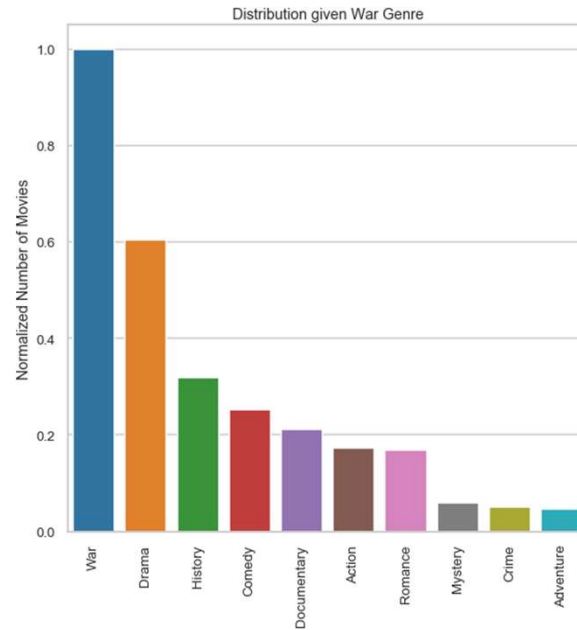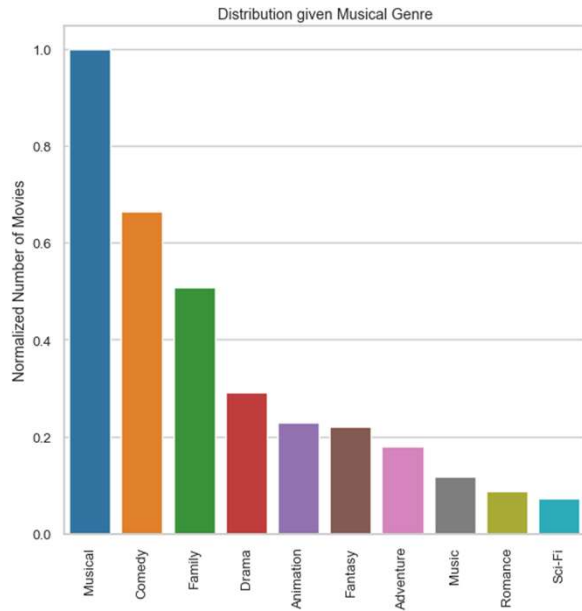- War & History

Genres with strong negative correlation

- Animation & Drama
- Comedy & Documentary

# Multi-Genre Distribution Plots

80% of the Romance, Mystery& Crime movies are also categorized as Drama

Multi-Genre Distribution Plots

65% of Musical movies are Comedy, 60% of War movies are Drama, Half of Short movies are Comedies
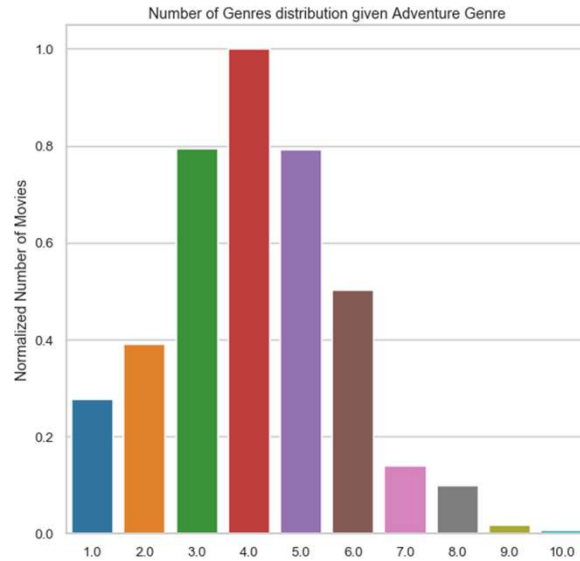
Number of Genres for [Sci-Fi, Adventure, Fantasy]

Movies categorized as Sci-Fi, Adventure or Fantasy are categorized into 3 to 6 categories
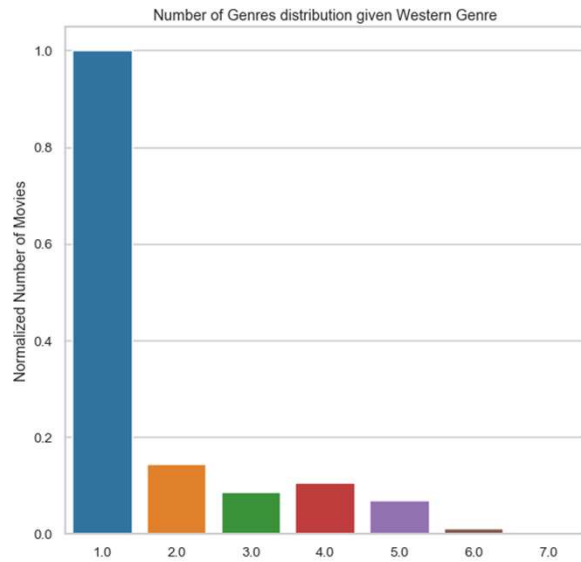
Number of Genres for [Western, Reality-TV, Adult]

Movie categorized as Western, Adult or Reality-TV are typically categorized to a single genre

Modeling

Modeling Overview

- Type: Supervised Learning

- Classification (Multi-label) Problem with 27 labels

- Pipeline consists of
  - Encoder
  - Classification Model

For all Pipeline (Encoder + Classification Model) Combinations

Pipeline Hyperparameter Tuning (via GridSearchCV)

Optimized Encoder + Model

Evaluate Encoder + Model (F1 Score)

Record Test F1 Score

Train Data

Test Data

Train/Test Split

1. Feature Engineering
2. Text Preprocessing

Input Data

Re-train Encoder + Model with entire Input Data

Best Encoder + Model Selection

Final Model

```
print(remove_tags('<html><h2>Learning NLP</h2></html>'))
print(remove_tags(' <a>Movie-Plot</a>'))
```
executed in 3ms, finished 15:08:11 2019-04-28

```
Learning NLP
Movie-Plot
```

Remove HTML Tags

Remove Punctuations

áàâäãå ÁÀÂÄÃÅ çÇ
éèêë ÉÈÊË íìîï ÍÌÎÏ
ñÑ óòôöõ ÓÒÔÖÕ
úùûü ÚÙÛÜ ÿŸ

Convert accented characters to ASCII

```
def keep_alpha(sentence):
    alpha_sentence = re.sub('[^a-z A-Z]+', ' ', sentence)
    return alpha_sentence
```

Keep only alphabetic strings

Text Preprocessing

Sample text with Stop Words	Without Stop Words
GeeksforGeeks – A Computer Science Portal for Geeks	GeeksforGeeks , Computer Science, Portal ,Geeks
Can listening be exhausting?	Listening, Exhausting
I like reading, so I read	Like, Reading, read

Stop words removal

	original_word	lemmatized_word
0	trouble	trouble
1	troubling	trouble
2	troubled	trouble
3	troubles	trouble

	original_word	lemmatized_word
0	goose	goose
1	geese	goose

Lemmatization

| Raw	Lowercased
Canada CanadA CANADA	canada

Lower casing all the words

Text Preprocessing

Train/Test Split

- Imbalanced Data with genre occurrences ranging from 61 (Adult) to 45891 (Drama)

- Train/Test split such that
  - At least 60% of the samples in the training set
  - At least 15% of the samples in the test set

Fraction of the number of movies of each Genre in Train and Test Sets
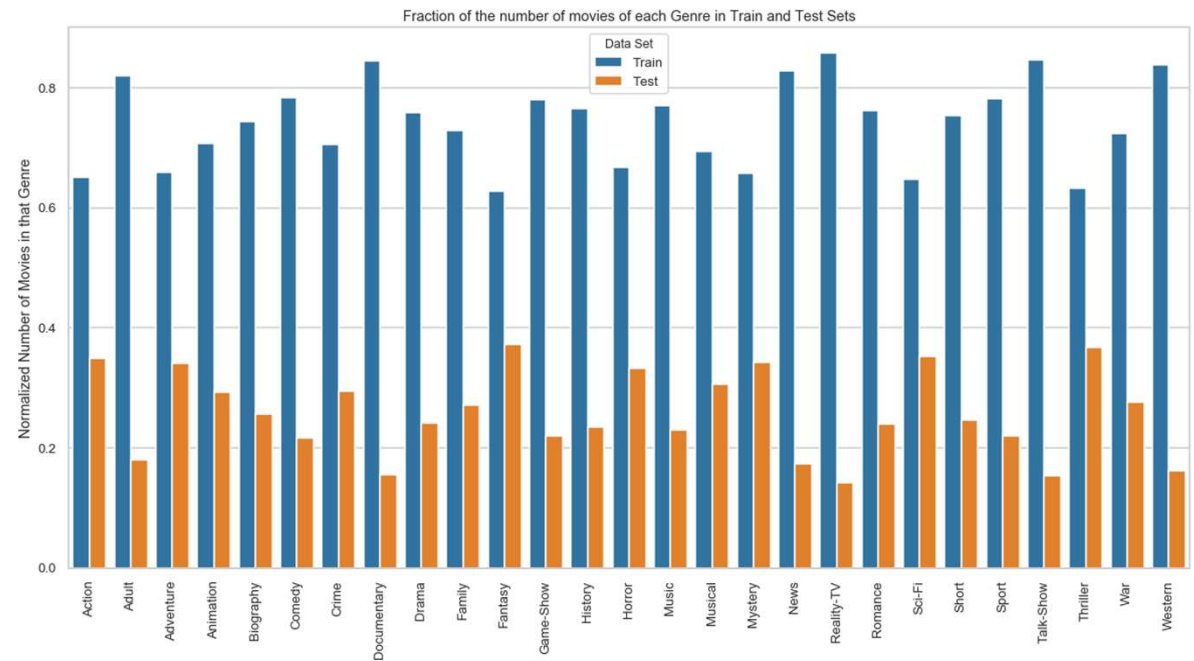
For each Genre

$$\text{Precision (Action)} = \frac{\text{Number of movies 'correctly' identified as Action Genre}}{\text{Total number of movies that have been identified as Action Genre}}$$

$$\text{Recall (Action)} = \frac{\text{Number of movies 'correctly' identified as Action Genre}}{\text{Total number of Action Genre movies in the data set}}$$

$$\text{F1 score (Action)} = \frac{2 * \text{Precision(Action)} * \text{Recall(Acation)}}{\text{Precision(Action)} + \text{Recall(Action)}}$$

Overall F1 Score = Weighted Average of individual Genre F1 Score

Evaluation Metric – Overall F1 Score

Multi-Label Algorithm
- Binary Relevance
- Label Powerset
- Label Powerset with Clustering

Text Encoder
- Count Vectorizer
- TF-IDF
- Sentence Embedding

Classifiers
- Logistic Regression, Linear SVC, Naïve Bayes
- Cosine Similarity
- Neural Networks

Multi-Label Algorithm + Text Encoder + Classifiers

Binary Relevance

- Treat each label as a separate single class Classification
- 27 Genres ➔ 27 Binary Classifiers

X	Y_1	Y_2	Y_3	Y_4
$x^{(1)}$	0	1	1	0
$x^{(2)}$	1	0	0	0
$x^{(3)}$	0	1	0	0
$x^{(4)}$	1	0	0	1
$x^{(5)}$	0	0	0	1

X	Y_1
$x^{(1)}$	0
$x^{(2)}$	1
$x^{(3)}$	0
$x^{(4)}$	1
$x^{(5)}$	0

X	Y_2
$x^{(1)}$	1
$x^{(2)}$	0
$x^{(3)}$	1
$x^{(4)}$	0
$x^{(5)}$	0

X	Y_3
$x^{(1)}$	1
$x^{(2)}$	0
$x^{(3)}$	0
$x^{(4)}$	0
$x^{(5)}$	0

X	Y_4
$x^{(1)}$	0
$x^{(2)}$	0
$x^{(3)}$	0
$x^{(4)}$	1
$x^{(5)}$	1

Label Powerset

- Treat each of the unique genre combinations found in the training data as a possible class
- 1505 Unique genre combinations ➔ Multi-class classification with 1505 classes

X	y1	y2	y3	y4
x1	0	1	1	0
x2	1	0	0	0
x3	0	1	0	0
x4	0	1	1	0
x5	1	1	1	1
x6	0	1	0	0

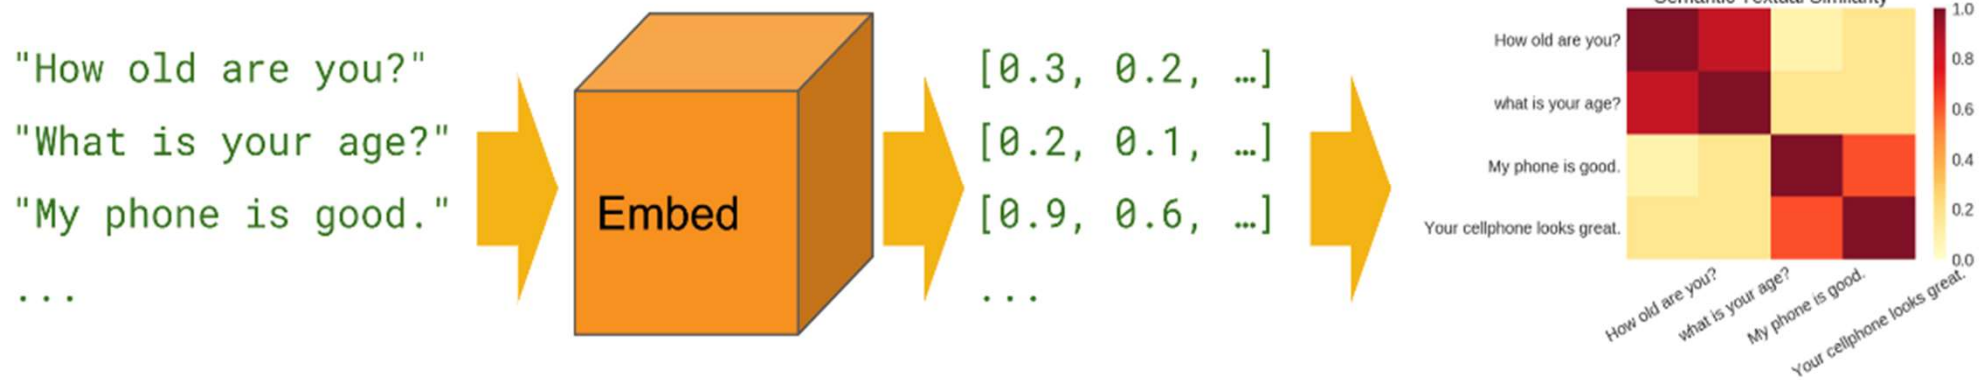| X	y1
x1	1
x2	2
x3	3
x4	1
x5	4
x6	3

Label Powerset with Clustering

- Reduce the number of genre combinations by clustering (from 1505 to 75)

# Multi-Label Classification Algorithms

Text Encoder – Sentence Embedding

- Google Universal Sentence Encoder (USE) converts input sentence to 512 dimension numeric vectors
- Preserves the semantic meaning of the sentence

Sentence Embedding – [Action, Horror, Game-Show]

Semantic Textual Similarity for: ['Action', 'Horror', 'Game-Show']

PCA Representation for Genres: ['Action', 'Horror', 'Game-Show']
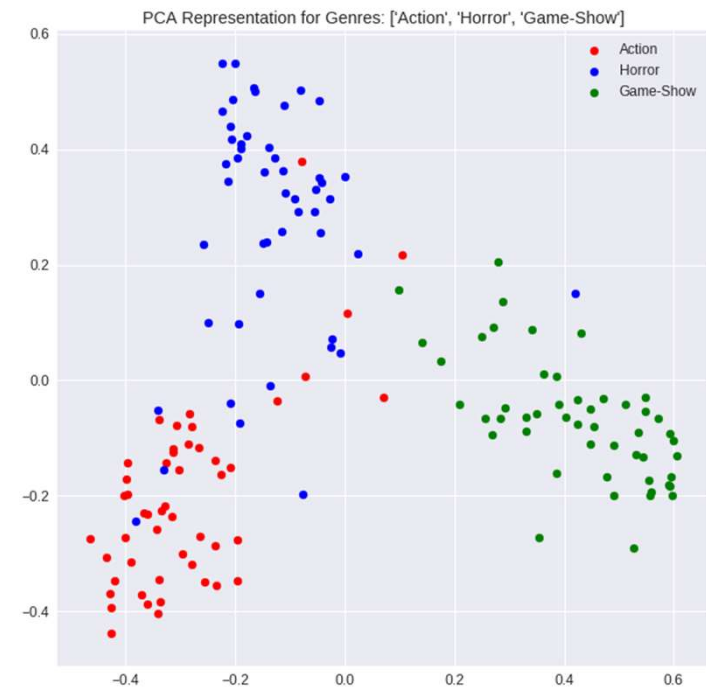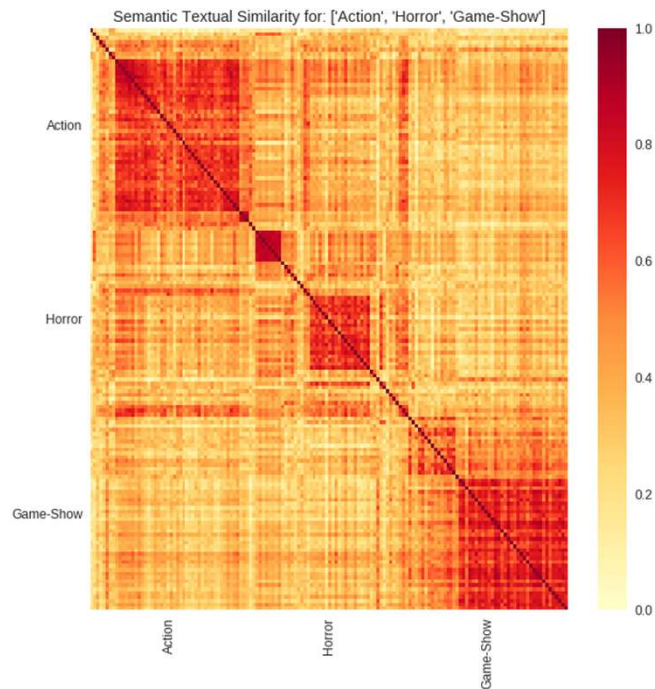
Sentence Embedding – [Adult, Crime, War]

Semantic Textual Similarity for: ['Adult', 'Crime', 'War']

PCA Representation for Genres: ['Adult', 'Crime', 'War']

Classification Models – Cosine Similarity

- Sentence Embedding of similar genres cluster together

- Cosine Similarity $(\vec{x}, \vec{y}) =$

- $\dfrac{\vec{x} \cdot \vec{y}}{||\vec{x}||\ ||\vec{y}||} = \dfrac{\sum_{i=1}^{512} x_i y_i}{\sqrt{\sum_{i=1}^{512} x_i^2}\ \sqrt{\sum_{i=1}^{512} y_i^2}}$

Y = Test Embedded Vector

$X_1 = [0.2, 0.1, ...-0.4]$ → Cosine Sim (X_1, Y)

$X_2 = [-0.2, 0.5, ...-0.3]$ → Cosine Sim (X_2, Y)

$X_{1505} = [0.32, -0.2, ...1.2]$ → Cosine Sim (X_{1505}, Y)

Argmax()= i → Y contains the genre combination as i Class
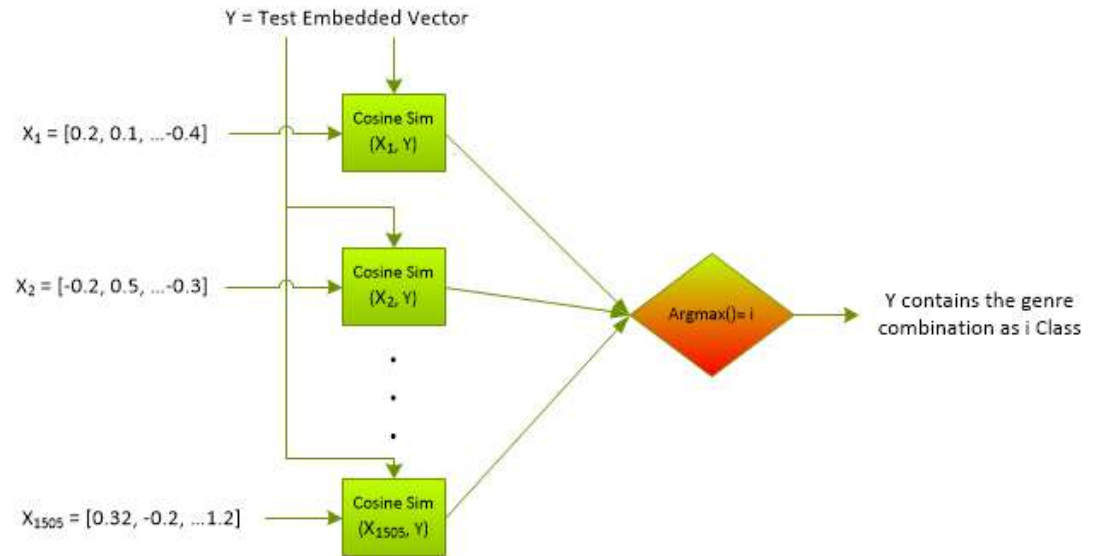
Classification Models – Neural Network

Input Layer: 512 nodes

Relu

Fully Connected with 800 nodes

Dropout (25%)

Sigmoid

Output Layer: 27 nodes

>thresh1? → Y → Genre1=1

>thresh2? → Y → Genre2=1

>thresh26? → Y → Genre26=1

>thresh27? → Y → Genre27=1
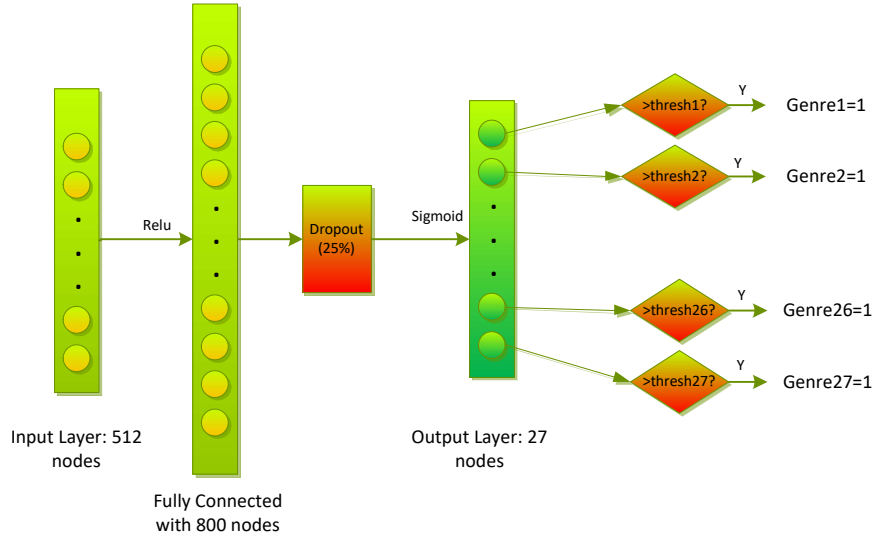
Using Binary Relevance
- Each Genre predicted separately ➔ Output Layer = 27 Neurons
- Use sigmoid Activation Function

Input Layer: 512 nodes

Relu

Fully Connected with 1024 nodes

Dropout (50%)

Softmax

Output Layer: 1505 nodes

Argmax()
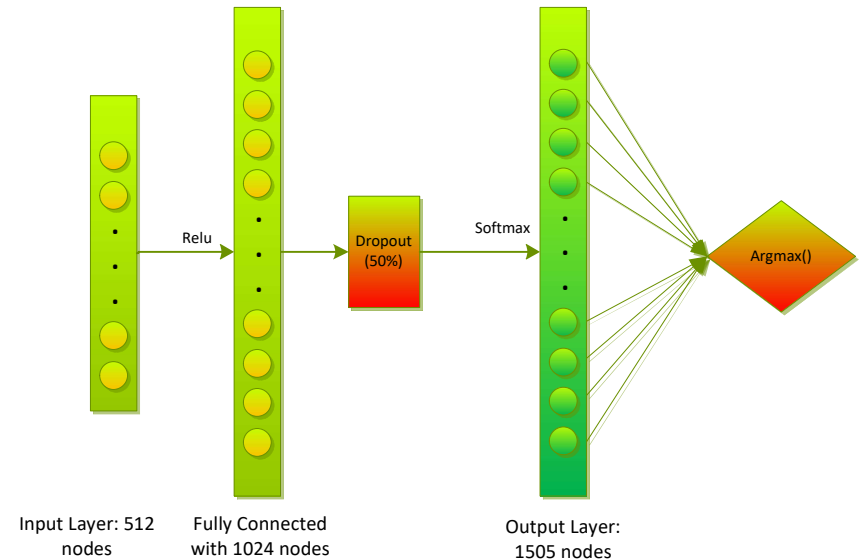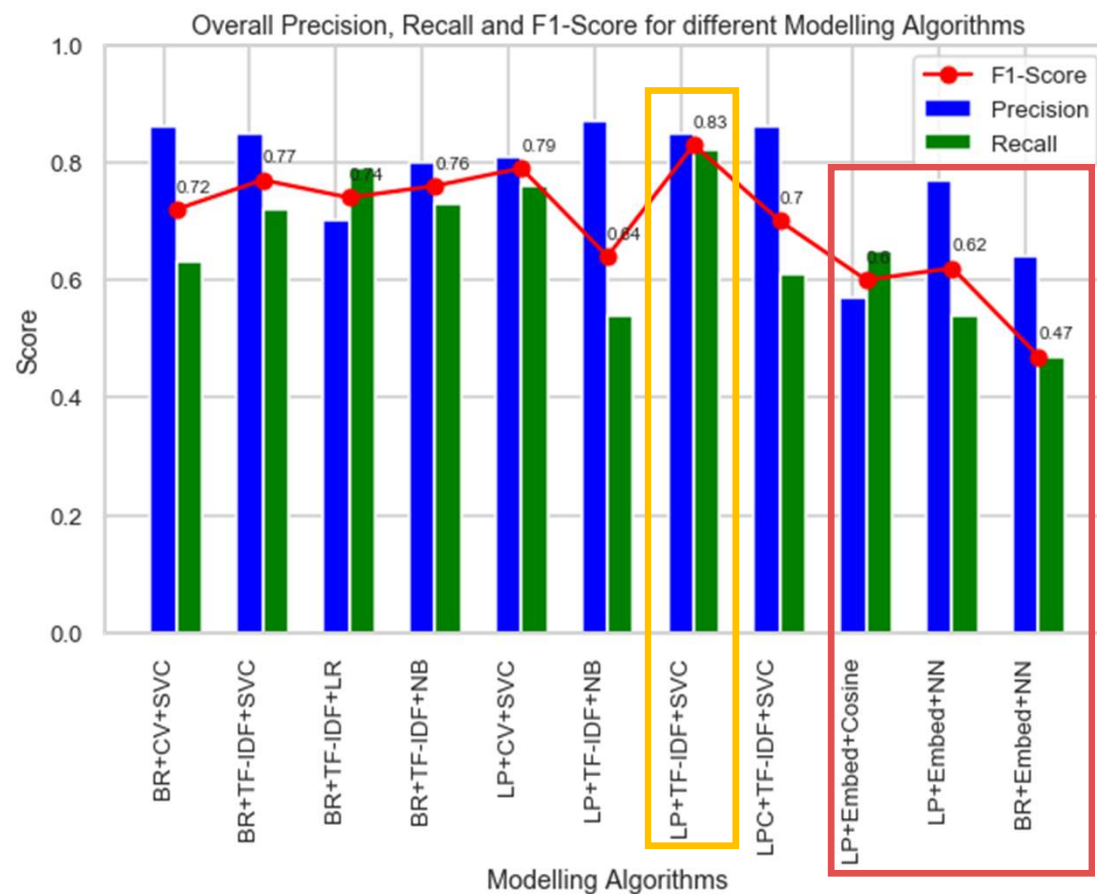
Using Label Powerset
- Output Layer = 1505 Neurons
- Use softmax Activation Function since a single class is predicted

Result Summary

- Best predicting Model: Label Powerset + TF-IDF + Linear SVC (Overall F1-Score = 0.83)

- Models using Sentence Embedding performance is worse compared to other Vectorizers and ML models

BR: Binary Relevance
LP: Label Powerset
LPC: Label Powerset with Clustering
CV: Count Vectorizer
Embed: Sentence Embedding via USE

SVC: Linear Support Vector Classifier
LR: Logistic Regression
NB: Naïve Bayes
Cosine: Cosine Similarity
NN: Neural Networks

Overall Precision, Recall and F1-Score for different Modelling Algorithms

Label Powerset + TF-IDF + Linear SVC

- Hyperparameters
  - TF-IDF
    - Ngram = (1, 2)
    - Min_df = 2
    - Max_df = 0.5
  - LinearSVC
    - C = 10
- Overall F1-score = 0.83
- Poorest performing Genre = Adult
  - F1-score = 0.32, but only 11 samples

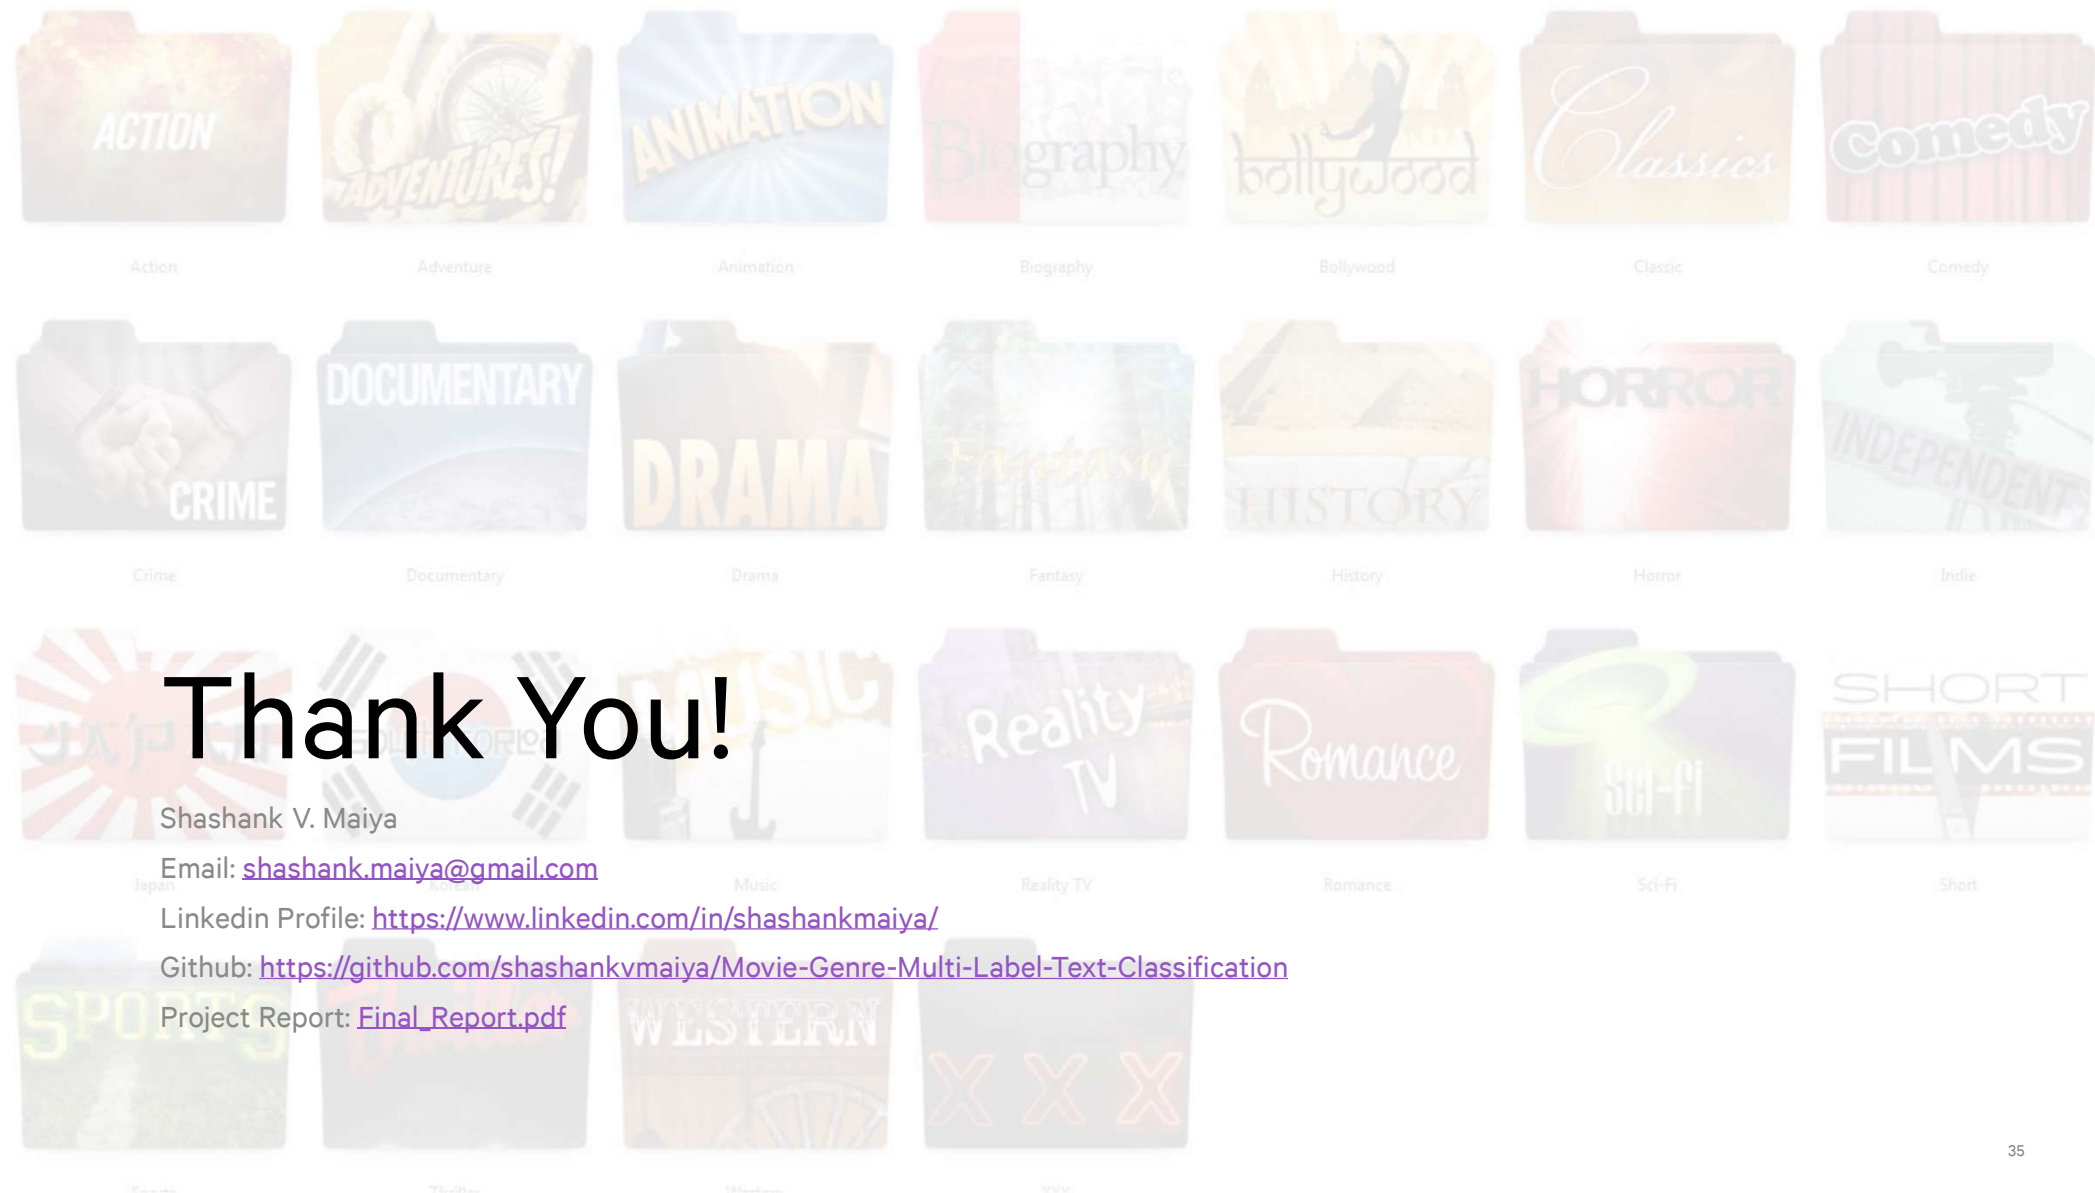|	Precision	Recall	F1-Score	Support
Action	0.87	0.82	0.84	4321.0
Adult	0.38	0.27	0.32	11.0
Adventure	0.85	0.80	0.82	3496.0
Animation	0.86	0.84	0.85	3333.0
Biography	0.54	0.43	0.48	354.0
Comedy	0.87	0.84	0.85	7320.0
Crime	0.86	0.86	0.86	4453.0
Documentary	0.72	0.73	0.73	1863.0
Drama	0.91	0.87	0.89	11067.0
Family	0.82	0.84	0.83	4173.0
Fantasy	0.86	0.78	0.81	2643.0
Game-Show	0.79	0.90	0.85	450.0
History	0.64	0.65	0.65	623.0
Horror	0.75	0.60	0.67	854.0
Music	0.76	0.80	0.78	654.0
Musical	0.75	0.65	0.70	182.0
Mystery	0.85	0.80	0.82	4114.0
News	0.76	0.80	0.78	681.0
Reality-TV	0.80	0.78	0.79	1748.0
Romance	0.87	0.86	0.87	4581.0
Sci-Fi	0.88	0.80	0.84	3055.0
Short	0.53	0.40	0.46	142.0
Sport	0.75	0.78	0.76	426.0
Talk-Show	0.77	0.86	0.81	809.0
Thriller	0.84	0.73	0.78	3254.0
War	0.74	0.71	0.72	388.0
Western	0.68	0.85	0.75	445.0
Avg/Total	0.85	0.82	0.83	65440.0

Conclusions

- Out of the 11 models considered, the best predicting model uses TF-IDF Vectorizer, Linear Support Vector Classifier and Label Powerset approach to achieve an overall F1-score of 0.83

- Sentence Embedding doesn't provide any benefit (yet)

- EDA specific Observations
    - Drama and Comedy the most popular Genre
    - On an average, a movie is classified into 2 genres (and a maximum of 12 genres)
    - Few strongly correlated genres include – a) Crime, Mystery & Thriller, b) Drama & Romance
    - 80% of Romance, Crime, Mystery and Thriller movies are also categorized as Drama

Limitations and Ideas

- Improving Sentence Embedding
  - Sentence embedding doesn't require the sentence lemmatization, or stop word removal, or in fact any of the text preprocessing steps. Use the original text before preprocessing to obtain sentence embedding
  - Use Sentence Embedding with Linear SVC (since Linear SVC was the best model)

Thank You!

Shashank V. Maiya

Email: shashank.maiya@gmail.com

Linkedin Profile: https://www.linkedin.com/in/shashankmaiya/

Github: https://github.com/shashankvmaiya/Movie-Genre-Multi-Label-Text-Classification

Project Report: Final_Report.pdf